

SNGULAR

LLM Distillation. Fine Tuning

Creating custom LLM model.

Contents:

- What is fine tuning(training-testing)
- Types of fine tuning
 - **Full Fine-Tuning**
 - **LoRA**
 - **QLoRA**
- Dataset preparation
 - **Vectorization**
 - **K-Means(Pairwise distance)**
- Hyperparameters
- Grid Search
- Fine tuning Use case
- Google Cloud Vertex AI.
- Decoder Strategies (Generation Parameters) for LLMs.
- Deployment & integration Use Case.

What is fine tuning about?

Base pretrained optimization function

Let:

- $\mathcal{D}_{\text{pretrain}}$: large dataset used for general pretraining.
- θ : the model parameters (weights) of the LLM.
- The goal is to minimize the loss:

$$\theta^* = \arg \min_{\theta} \mathbb{E}_{(x,y) \sim \mathcal{D}_{\text{pretrain}}} [\mathcal{L}(f(x; \theta), y)]$$

where:

- x : input (e.g. tokenized text),
- y : target (e.g. next token in causal language modeling),
- \mathcal{L} : loss function (e.g. cross-entropy),
- $f(x; \theta)$: output logits of the model.

Gradient update rule

$$\theta^{(t+1)} = \theta^{(t)} - \eta \cdot \nabla_{\theta} \mathcal{L}(f(x; \theta^{(t)}), y)$$

where:

- η : learning rate,
- ∇_{θ} : gradient with respect to the model parameters.

Fine tuning optimization function

$$\theta_{\text{fine}} = \arg \min_{\theta} \mathbb{E}_{(x,y) \sim \mathcal{D}_{\text{fine}}} [\mathcal{L}(\text{DistilledLLM}(x; \theta), y)] \quad \text{with} \quad \theta \leftarrow \theta_S^*$$

Learning Process

Train

- This is the **data the model learns from**.
- It **includes both** the input features (**Symptoms**) and the correct output (**Diagnostic**).
- The model uses this data to find **patterns or relationships**.

Test

- This is **new data the model hasn't seen before**.
- It's used to **check how well the model performs** on unseen examples.
- It helps evaluate whether the model **generalizes well**.

Example

Training a model to recognize cats vs. dogs in pictures:

- **Training Set:** 80 photos (40 cats, 40 dogs) → Model learns features (like ears, tails).
- **Test Set:** 20 new photos → We check if the model can still correctly tell cats from dogs.

Full Fine Tuning

Cross-Entropy Loss

$$\mathcal{L}_{\text{CE}} = - \sum_{t=1}^T \log P_{\theta}(y_t \mid x_{<t})$$

Gradient Descent or Adam

Updates all parameters θ

$$\theta \leftarrow \theta - \eta \cdot \nabla_{\theta} \mathcal{L}(f(x; \theta), y)$$

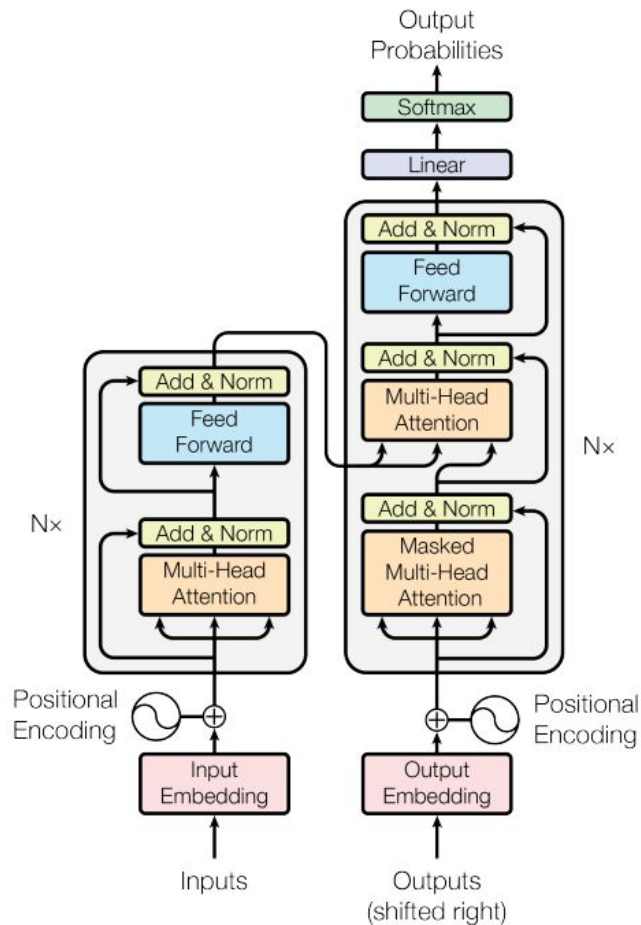



Figure 1: The Transformer - model architecture.

LoRA Fine Tuning


$$W' = W + \Delta W \quad \text{where } \Delta W = BA, \quad A \in \mathbb{R}^{r \times k}, \quad B \in \mathbb{R}^{d \times r}$$

$$h = Wx + BAx$$

Scaled Dot-Product Attention



Multi-Head Attention



QLoRA Fine Tuning

Instead of using:

$$h = Wx$$

QLoRA uses:

$$h = \underbrace{\tilde{W}x}_{\text{quantized frozen weights}} + \underbrace{\alpha \cdot BAx}_{\text{trainable LoRA adapters}}$$

$$\min_{\phi} \frac{1}{N} \sum_{i=1}^N \mathcal{L}(f(x_i; \tilde{\theta}, \phi), y_i)$$

- $\tilde{\theta}$: quantized model weights (frozen)
- ϕ : LoRA adapters (trainable)
- Only ϕ gets updated during training

- \tilde{W} is stored in **4-bit quantized form** (typically using `NF4` quantization).
- The adapter BA is **full-precision (float32 or bfloat16)**.

Dataset Preparation

Vectorization

1- Input Sentence

"Hello"

2- Character Vocabulary

Assume a vocabulary:

$\mathcal{V} = \{H, e, l, o\} \Rightarrow$ Indices: $H = 0, e = 1, l = 2, o = 3$

3- One-Hot Encoding

Each character is converted into a **one-hot vector** of size 4 (the vocabulary size):

Character	One-hot vector
H	[1, 0, 0, 0]
e	[0, 1, 0, 0]
l	[0, 0, 1, 0]
o	[0, 0, 0, 1]

So the word "Hello" becomes:

$$X = \begin{bmatrix} [1, 0, 0, 0] \\ [0, 1, 0, 0] \\ [0, 0, 1, 0] \\ [0, 0, 1, 0] \\ [0, 0, 0, 1] \end{bmatrix} \in \mathbb{R}^{5 \times 4}$$

Dataset Preparation

Vectorization

4- Embedding Matrix

Let's define a **learned embedding matrix** $E \in \mathbb{R}^{4 \times 3}$ to map each character to a 3D vector:

$$E = \begin{bmatrix} \text{H} & \rightarrow & [0.1, 0.3, 0.5] \\ \text{e} & \rightarrow & [0.2, 0.1, 0.4] \\ \text{l} & \rightarrow & [0.4, 0.4, 0.4] \\ \text{o} & \rightarrow & [0.7, 0.9, 0.2] \end{bmatrix}$$

$$X \cdot E = Z \in \mathbb{R}^{5 \times 3}$$

$$Z = \begin{bmatrix} [0.1, 0.3, 0.5] \\ [0.2, 0.1, 0.4] \\ [0.4, 0.4, 0.4] \\ [0.4, 0.4, 0.4] \\ [0.7, 0.9, 0.2] \end{bmatrix} \in \mathbb{R}^{5 \times 3}$$

Dataset Preparation


K-Means: Statistically Representative Sample

$$\min_{\{C_1, \dots, C_K\}} \sum_{k=1}^K \sum_{x_i \in C_k} \|x_i - \mu_k\|^2$$

Where:

- C_k : the set of points assigned to cluster k
- μ_k : the **centroid** (mean) of cluster C_k :

$$\mu_k = \frac{1}{|C_k|} \sum_{x_i \in C_k} x_i$$



Hyperparameters


Learning Rate

For a model with parameters θ , and a loss function $\mathcal{L}(\theta)$, the gradient update is:

$$\theta^{(t+1)} = \theta^{(t)} - \eta \cdot \nabla_{\theta} \mathcal{L}(\theta^{(t)})$$

Where:

- $\theta^{(t)}$: parameters at iteration t
- $\nabla_{\theta} \mathcal{L}$: gradient of the loss w.r.t. parameters
- η : learning rate



Hyperparameters


Batch Size

Stochastic Gradient Descent is approximated to the true gradient:

$$\nabla_{\theta} \mathcal{L}_{\text{true}}(\theta) = \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} \mathcal{L}(\theta; x_i, y_i)$$

With the **mini-batch gradient**:

$$\nabla_{\theta} \mathcal{L}_{\text{batch}}(\theta) = \frac{1}{B} \sum_{(x_j, y_j) \in \mathcal{B}} \nabla_{\theta} \mathcal{L}(\theta; x_j, y_j)$$



Hyperparameters


Gradient Accumulation

- b : micro-batch size (e.g., 4)
- g : gradient accumulation steps

$$B = b \cdot g$$



$$\nabla_{\theta} \mathcal{L}_B = \frac{1}{B} \sum_{i=1}^B \nabla_{\theta} \mathcal{L}(\theta; x_i)$$

$$\theta \leftarrow \theta - \eta \cdot \nabla_{\theta} \mathcal{L}_B$$



Hyperparameters

Grid Search



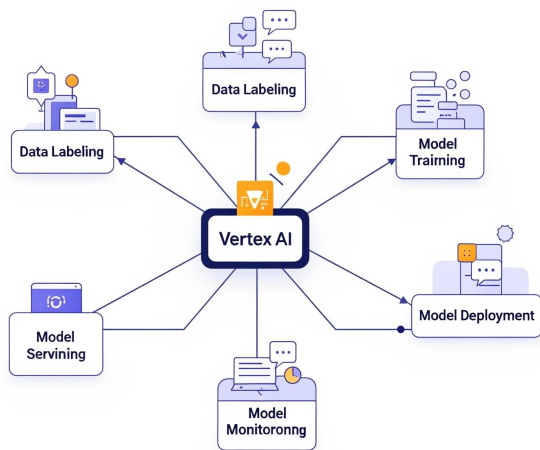
Use Case

- 01** Find Best Hyperparameters
- 02** Fine Tune Qwen Distilled Model
- 03** Push to Hugging Face



Deployment & Integration

LLM Deployment And Service Integration



What is Vertex AI?

- Managed Machine Learning Platform
- Supports Custom and AutoML Models.
- Vertex AI accelerates MLOps workflows efficiently.
- It integrates seamlessly with other Google Cloud services.
- Jupyter notebooks and Training Infrastructure.
- Scalable and Cost-Effective(Autopilot).

Python SDK Google Cloud Vertex AI

- The SDK is part of the `google-cloud-aiplatform` package.
- Simplifies working with models, datasets, endpoints, pipelines, and jobs.
- Initialization:

```
from google.cloud import aiplatform

aiplatform.init(
    project = "your-project-id",
    location = "us-central1"
)
```

- Create endpoint reference:

```
ENDPOINT = aiplatform.Endpoint(
    endpoint_name=(
        f"projects/your-project-id"
        f"/locations/us-central1"
        f"/endpoints/your-endpoint-id"
    )
)
```

- Make prediction:

```
prediction = ENDPOINT.predict(
    instances=["".join(prompt)],
    parameters={
        "temperature": 1.0,
        "max_new_tokens": 512,
        "top_k": 50,
        "top_p": 1.0,
        "repetition_penalty": 1.0
    }
)
```



Cloud SDK
CLI for GCP

Generation Parameters

Temperature (randomness of predictions)

Let:

- z_i be the logit (unnormalized score) for token i
- V be the vocabulary size

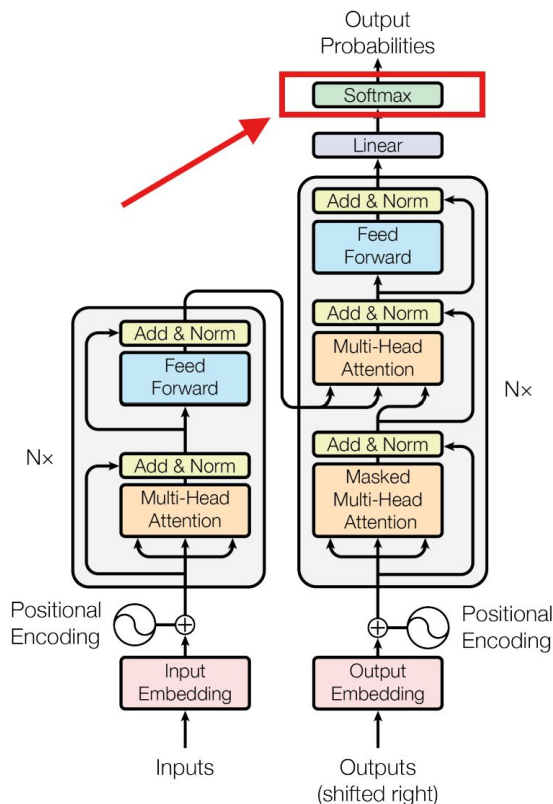
The **standard softmax** computes the probability p_i of token i as:

$$p_i = \frac{e^{z_i}}{\sum_{j=1}^V e^{z_j}}$$

To apply temperature $T > 0$, the logits are **divided by T** :

$$p_i^{(T)} = \frac{e^{z_i/T}}{\sum_{j=1}^V e^{z_j/T}}$$

- $T = 1$: Regular softmax (default).
- $T < 1$: Makes the distribution **sharper** (more confident).
- $T > 1$: Makes the distribution **flatter** (more random).



Generation Parameters

Top K (pick top-k most likely tokens)

Let:

- z_i : logit for token i
- V : vocabulary of size N
- p_i : softmax probability for token i

Compute probabilities:

$$p_i = \frac{e^{z_i}}{\sum_{j=1}^N e^{z_j}}$$

- Identify the set $S_k \subseteq V$ of the k tokens with the highest probabilities.

$$S_k = \{i \in V \mid p_i \text{ is among the top } k \text{ values in } \{p_j\}_{j=1}^N\}$$

Example (k = 3)

Token	p_i
A	0.50
B	0.25
C	0.15
D	0.07
E	0.03

Generation Parameters

Top P (pick group of tokens with total probability $\geq p$)

Let:

- z_i : logit for token $i \in V$, the vocabulary
- p_i : softmax probability

$$p_i = \frac{e^{z_i}}{\sum_{j=1}^N e^{z_j}}$$

Create a sorted list of tokens $T = \{t_1, t_2, \dots, t_N\}$, such that:

$$p_{t_1} \geq p_{t_2} \geq \dots \geq p_{t_N}$$

Define S_p as the smallest set of tokens such that:

$$\text{cumulative probability } \sum_{t_i \in S_p} p_{t_i} \geq p$$

Example (top-p = 0.9)

Token	p_i
A	0.40
B	0.30
C	0.15
D	0.10
E	0.05

Cumulative sum:

- A: 0.40
- A + B: 0.70
- A + B + C: 0.85
- A + B + C + D: **0.95** → first time ≥ 0.9

Generation Parameters

Repetition Penalty (penalizes frequent tokens to avoid repetition)

Let:

- \mathcal{V} : vocabulary
- $L \in \mathbb{R}^{|\mathcal{V}|}$: the vector of logits output by the model at time step t
- $R > 1$: the **repetition penalty** factor (e.g., 1.1, 1.2, ...)

Let $G = \{g_1, g_2, \dots, g_{t-1}\}$ be the set of **previously generated tokens**.

For each token $i \in \mathcal{V}$:

$$L'_i = \begin{cases} \frac{L_i}{R} & \text{if } i \in G \text{ and } L_i > 0 \\ L_i \cdot R & \text{if } i \in G \text{ and } L_i < 0 \\ L_i & \text{otherwise} \end{cases}$$

$$p_i = \frac{e^{L'_i}}{\sum_{j \in \mathcal{V}} e^{L'_j}}$$

Example ($R = 1.2$)

- Vocabulary: ["cat", "dog", "mouse"]
- Original logits: $L = [3.0, 1.0, 0.5]$
- Generated so far: ["cat"]

Apply the penalty:

- "cat": $3.0 / 1.2 = 2.5$
- "dog" and "mouse" stay the same

New logits: $[2.5, 1.0, 0.5]$ → lower chance of picking "cat" again

Bibliography

- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... & Polosukhin, I. (2017). Attention is all you need. *Advances in neural information processing systems*, 30.
- Cho, K., Van Merriënboer, B., Gulcehre, C., Bahdanau, D., Bougares, F., Schwenk, H., & Bengio, Y. (2014). Learning phrase representations using RNN encoder-decoder for statistical machine translation. *arXiv preprint arXiv:1406.1078*.
- Hu, E. J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., ... & Chen, W. (2022). Lora: Low-rank adaptation of large language models. *ICLR*, 1(2), 3.
- Dettmers, T., Pagnoni, A., Holtzman, A., & Zettlemoyer, L. (2023). Qlora: Efficient finetuning of quantized llms. *Advances in neural information processing systems*, 36, 10088-10115.
- Hamerly, G., & Elkan, C. (2003). Learning the k in k-means. *Advances in neural information processing systems*, 16.
- Zhu, Y., Li, J., Li, G., Zhao, Y., Jin, Z., & Mei, H. (2024, March). Hot or cold? adaptive temperature sampling for code generation with large language models. In Proceedings of the AAAI Conference on Artificial Intelligence (Vol. 38, No. 1, pp. 437-445).
- Yerram, V., You, C., Bhojanapalli, S., Kumar, S., Jain, P., & Netrapalli, P. (2024). HiRE: High Recall Approximate Top-\$ k \$ Estimation for Efficient LLM Inference. *arXiv preprint arXiv:2402.09360*.
- Chu, K., Chen, Y. P., & Nakayama, H. (2024). A better llm evaluator for text generation: The impact of prompt output sequencing and optimization. *arXiv preprint arXiv:2406.09972*.

Next Talk

- Retrieval-augmented Generation (RAG).
- Grounding and Context.
- Vectorized Database - Distance Measurements.
- Real-world use case:
 - RAG implementation to enhance community well-being.

Questions

Thanks for your attention!!!